



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

Python for IoT Data Analytics

Time-series Forecasting

EIT4SEMM PHD Program, July 2023

Prof. Marco Di Felice

Department of Computer Science and Engineering, University of Bologna

marco.difelice3@unibo.it



IoT Sensor Data Forecasting

For ease of disposition, we distinguish between:

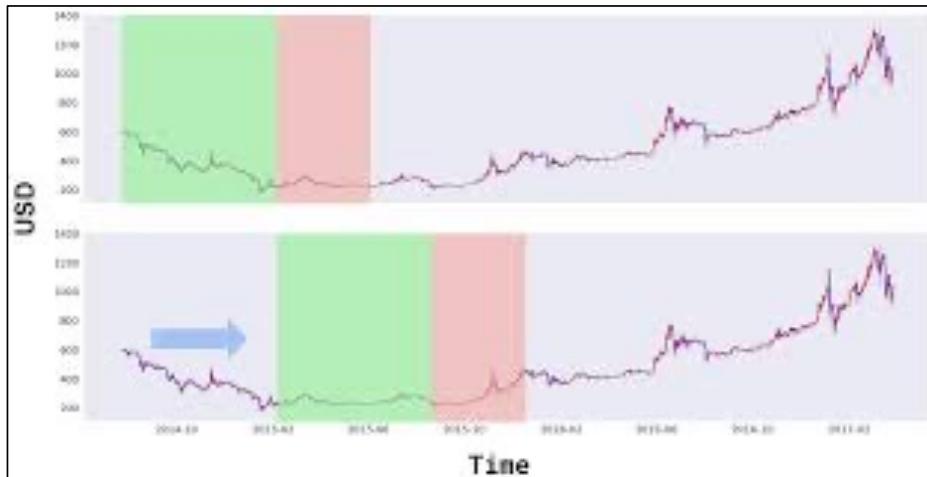
- ✧ **Time-series Processing** → Techniques and tools to **manipulate** the time-series (e.g. data aggregation).
- ✧ **Time-series Analytics** → Techniques and tools to **extract information** from the time-series (e.g. forecasting).

NOTE. In several use-cases, processing and analytics are combined together (e.g. first aggregate/filter the time-series, then apply learning techniques).



IoT Sensor Data Forecasting

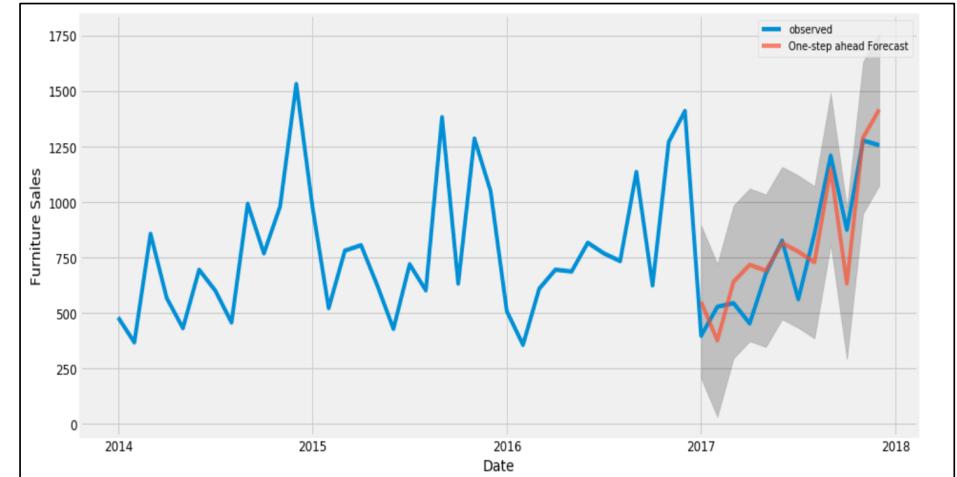
Source: <https://www.mdpi.com/1911-8074/12/1/17/htm>



TIME-SERIES CLASSIFICATION

- Assign a label to a time-series or to a subset

Source: <https://towardsdatascience.com/an-end-to-end-project-on-time-series-analysis-and-forecasting-with-python-4835e6bf050b>



TIME-SERIES FORECASTING

- Predict future values based on previously observed values.



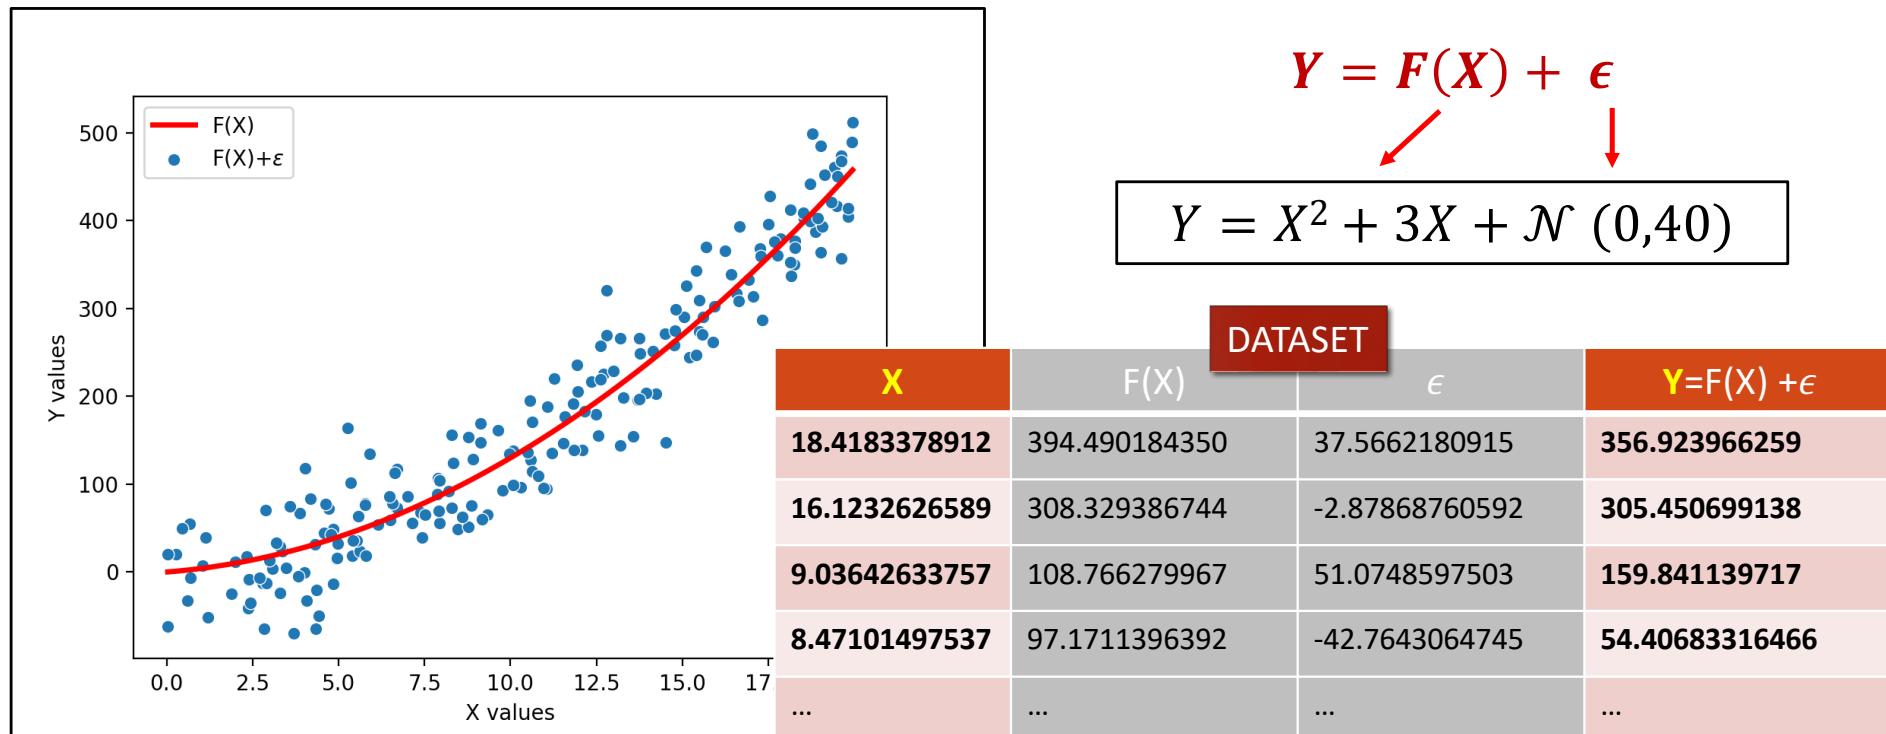
IoT Sensor Data Forecasting

PREDICTION/FORECASTING

- ❑ Given a quantitative response y and p different variables (x_1, x_2, \dots, x_p)
- ❑ Find the relationship between y and $\mathbf{x} = (x_1, x_2, \dots, x_p)$: $y = f(\mathbf{x}) + \epsilon$
- ❑ ϵ is a random error term, which is independent of \mathbf{x} and has mean zero.
- ❑ The function f that connects the input variable to the output variable is **unknown**
- ❑ Learn the function $\hat{y} = \hat{f}(\mathbf{x})$, where \hat{f} represents our estimate for f , and \hat{y} represents the resulting prediction for y



IoT Sensor Data Forecasting





IoT Sensor Data Forecasting

CLASSIFICATION

- ❑ Given a **qualitative** response y and p different variables (x_1, x_2, \dots, x_p)
- ❑ The variable y can take values in one of K different **classes** or **categories**
- ❑ Like in prediction problems, the function f that connects the input variable to the output variables is **unknown**
- ❑ Learn the function $\hat{y} = \hat{f}(x)$ with $\hat{f} : X \rightarrow K$ **minimizing the error rate**:

$$\frac{1}{n} \sum_{i=1}^n I_i \quad I_i = \begin{cases} 1 & \text{if } (y_i \neq \hat{y}_i) \\ 0 & \text{if } (y_i = \hat{y}_i) \end{cases}$$



IoT Sensor Data Forecasting

Time series Components

✧ Secular Trend

- Smooth long-term direction of a time series

✧ Seasonal

- Patterns of change in a time series within a period. These patterns tend to repeat themselves each period.

✧ Irregular Variations

- Further classified as episodic or residual, cannot be predicted



IoT Sensor Data Forecasting

Single-variate analysis

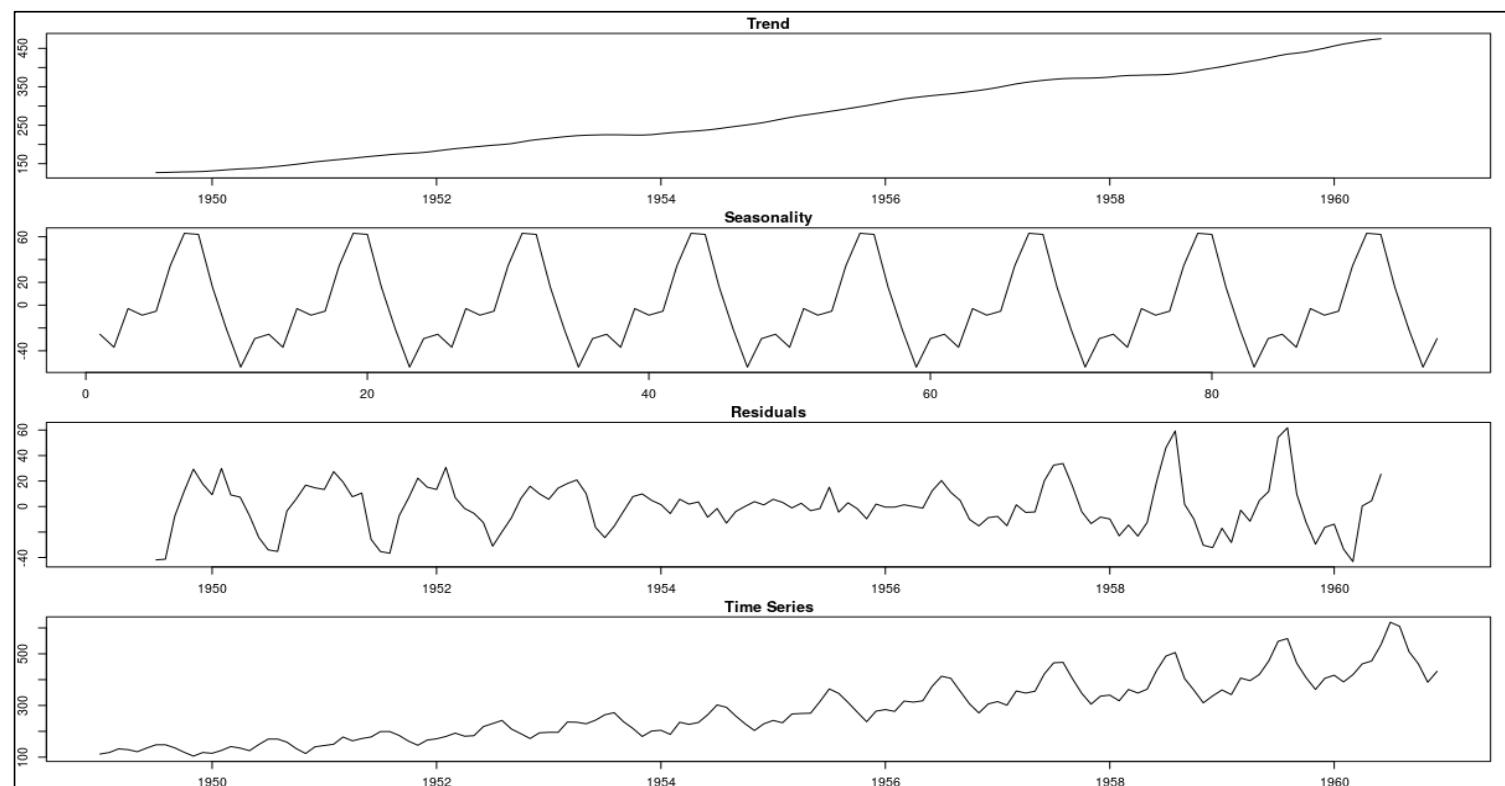
Date	Soil Moisture
2023-07-20 08:00:00	34.12
2023-07-20 09:00:00	38.12
2023-07-20 10:00:00	37.12
2023-07-20 11:00:00	39.12
2023-07-20 12:00:00	39.45
2023-07-20 13:00:00	60.12

Multi-variate analysis

Date	Soil Moisture	Precipitation
2023-07-20 08:00:00	34.12	0
2023-07-20 09:00:00	38.12	0
2023-07-20 10:00:00	37.12	0
2023-07-20 11:00:00	39.12	0
2023-07-20 12:00:00	39.45	0
2023-07-20 13:00:00	60.12	2



IoT Sensor Data Forecasting



TREND

SEASONALITY

RESIDUALS

TIME-SERIES



IoT Sensor Data Forecasting

ADDITIVE DECOMPOSITION

Observed Series = Trend + Seasonal + Irregular

$$O_t = T_t + S_t + I_t$$

MULTIPLICATIVE DECOMPOSITION

Observed Series = Trend * Seasonal * Irregular

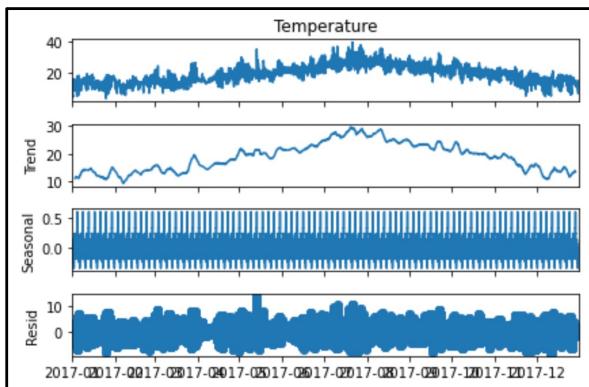
$$O_t = T_t * S_t * I_t$$



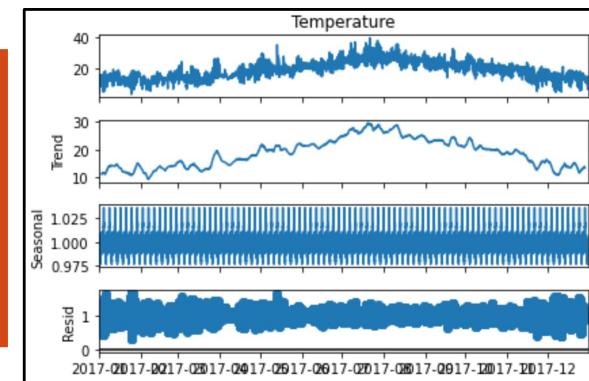
IoT Sensor Data Forecasting

```
from statsmodels.tsa.seasonal import seasonal_decompose  
result = seasonal_decompose(df["Temperature"],  
                           model="multiplicative",  
                           period = 60*10)  
result.plot()
```

ADDITIVE

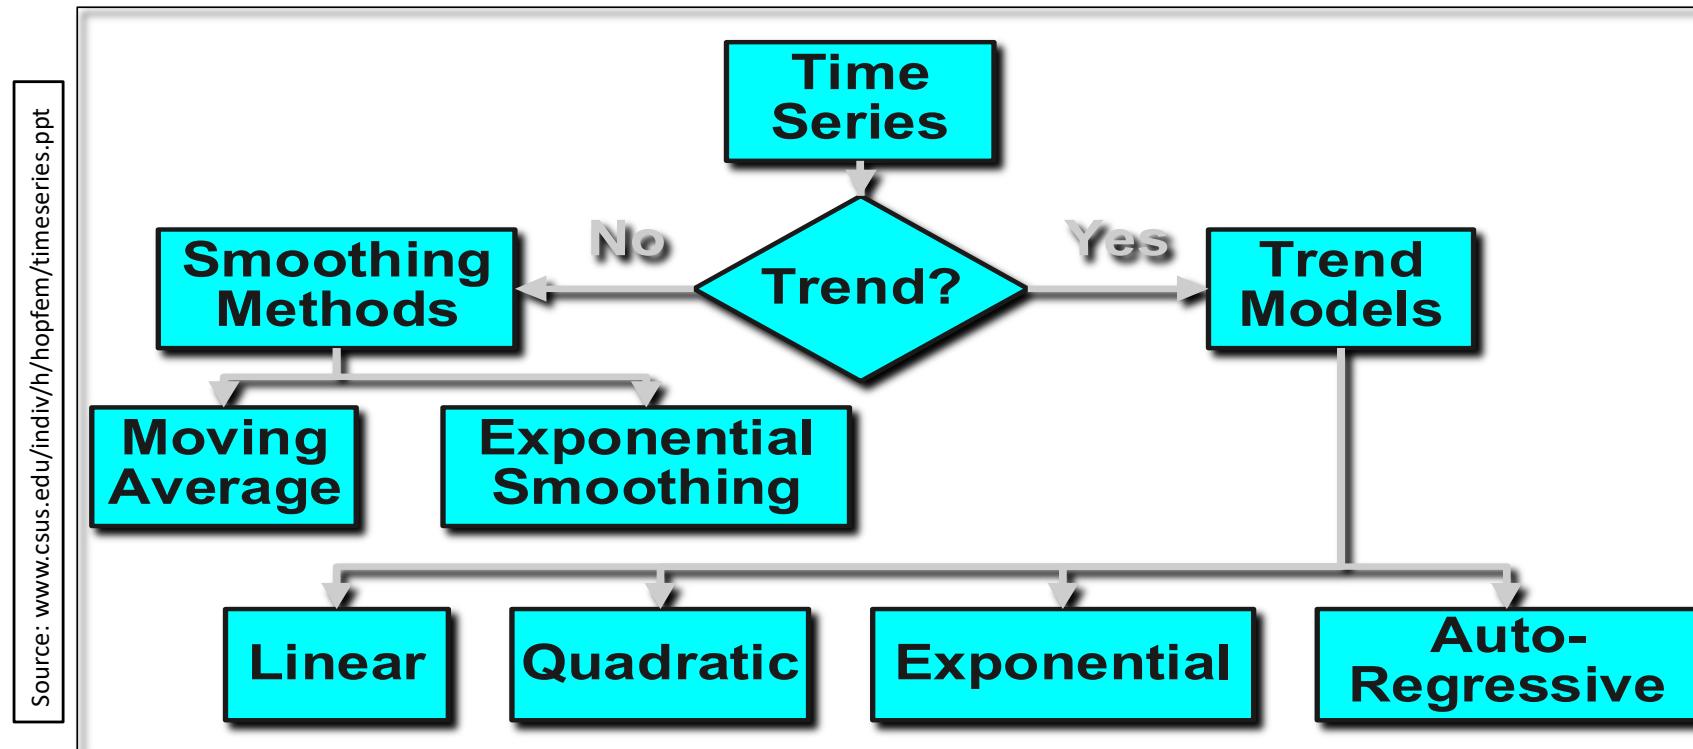


MULTIPLICATIVE





IoT Sensor Data Forecasting





IoT Sensor Data Forecasting

(SIMPLE) SMOOTHING
METHODS



Source: www.csus.edu/indiv/h/hopfem/timeseries.ppt

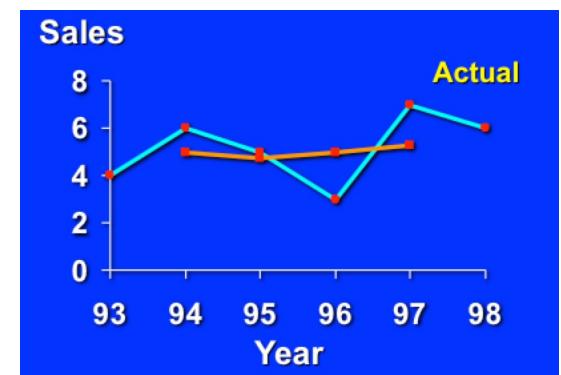


IoT Sensor Data Forecasting

✧ Moving Average Method

- Elementary **forecasting** solution
- The n -period forecast is simply **the average of the observations** in the most recent n periods:

$$A_t = (x_t + x_{t-1} + \dots + x_{t-n+1}) / n$$





Source: www.csus.edu/indiv/h/hopfem/timeseries.ppt

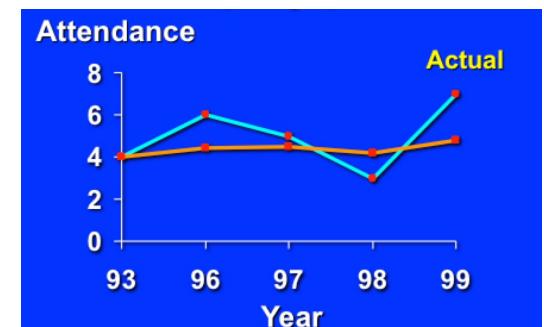


IoT Sensor Data Forecasting

✧ Exponential Smoothing

- Form of weighted moving average
- Use a **smoothing costant** (W), range between 0 and 1
- Record only **actual value** and **past prediction**

$$S_t = \alpha x_t + (1 - \alpha)S_{t-1}$$





IoT Sensor Data Forecasting

LINEAR REGRESSION METHODS



Source: R.B. Clough – UNH M. E. Henrie - UAA



IoT Sensor Data Forecasting

- ❑ Straightforward approach for predicting a quantitative response
- ❑ Parametric method, *explainable* modeling
- ❑ Assume a *linear relationship* between X and Y :

LINEAR REGRESSION

$$\hat{y} = \hat{f}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p$$



Estimate $\hat{f}(x)$ → Estimate $\beta_0, \beta_1, \beta_2 \dots \beta_p$



Use the **training data** to fit/train the model ...



Source: R.B. Clough – UNH M. E. Henrie - UAA



IoT Sensor Data Forecasting

□ Simple Linear Regression ($p=1$)

Dataset composed of n observation pairs:

$$(x_1, y_1), (x_2, y_2), \dots (x_n, y_n)$$

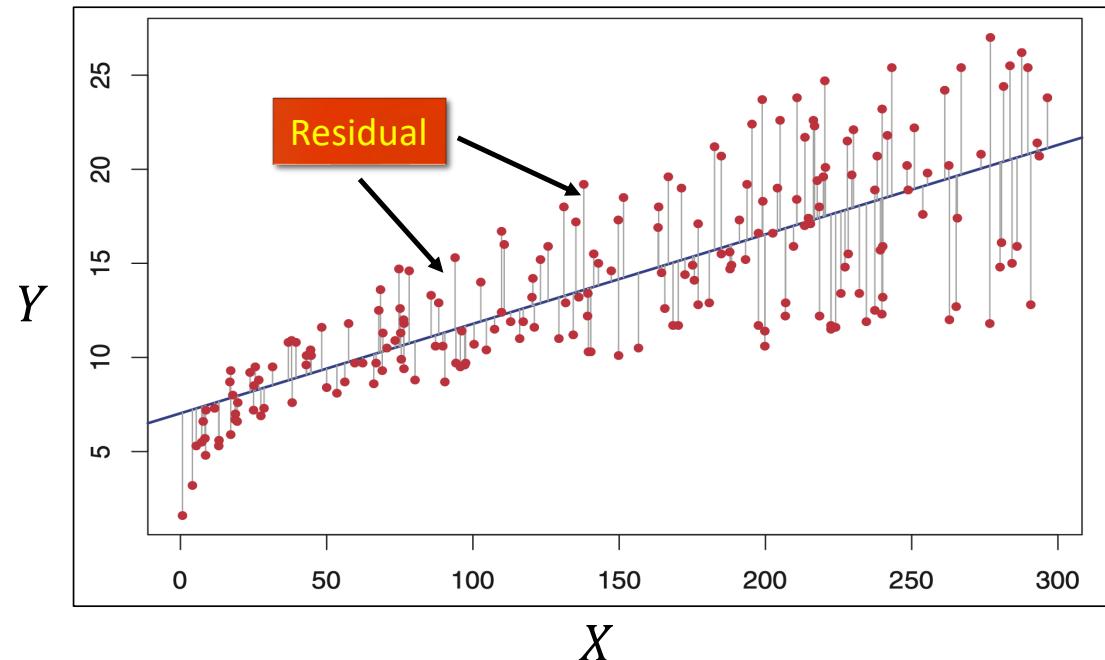
Regression problem:

$$\hat{y} = \hat{f}(\mathbf{x}) = \beta_0 + \beta_1 x$$



IoT Sensor Data Forecasting

□ Simple Linear Regression ($p=1$)



Real Y values:

$$y_1, y_2, y_3, \dots, y_n$$

Predicted $\hat{Y} = \hat{f}(x)$ values:

$$\hat{y}_1, \hat{y}_2, \hat{y}_3, \dots, \hat{y}_n$$

Residual of the i -th prediction:

$$e_i = y_i - \hat{y}_i$$

Residual Sum of Squares (RSS):

$$\sum_{i=1}^n e_i^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



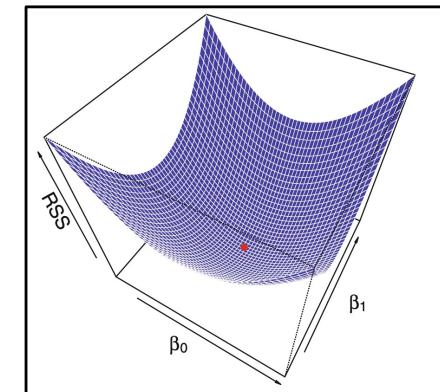
IoT Sensor Data Forecasting

□ Simple Linear Regression ($p=1$)

LINEAR LEAST SQUARE METHODS

Compute β_0, β_1 values such that the **RSS function is minimized**

$$\begin{aligned} \operatorname{argmin}_{\beta_0, \beta_1} \text{RSS} &= \sum_1^n (y_i - \hat{y}_i)^2 \\ &= \sum_1^n (y_i - \beta_0 - \beta_1 x_i)^2 \end{aligned}$$





IoT Sensor Data Forecasting

□ Simple Linear Regression ($p=1$)

LINEAR LEAST SQUARE METHODS

$$RSS = \sum_{1}^n (y_i - \beta_0 - \beta_1 x_i)^2$$

$$\begin{cases} \frac{\partial RSS}{\partial \beta_0} = -2 \sum_{1}^n (y_i - \beta_0 - \beta_1 x_i) = 0 \\ \frac{\partial RSS}{\partial \beta_1} = -2 x_i \sum_{1}^n (y_i - \beta_0 - \beta_1 x_i) = 0 \end{cases}$$



IoT Sensor Data Forecasting

□ Linear Regression

$$\hat{y} = \hat{f}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p$$

Rewriting the system of linear Equations with **matrices**:

$$Y = X\beta$$
$$Y = \begin{pmatrix} y_1 \\ \dots \\ y_n \end{pmatrix}$$
$$X = \begin{pmatrix} 1 & x_{1,1} & \dots & x_{1,p} \\ 1 & x_{1,\dots} & \dots & \dots \\ 1 & x_{1,1} & \dots & x_{n,p} \end{pmatrix}$$
$$\beta = \begin{pmatrix} \beta_0 \\ \dots \\ \beta_p \end{pmatrix}$$



IoT Sensor Data Forecasting

□ Linear Regression

$$\hat{y} = \hat{f}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p$$

LINEAR LEAST SQUARE METHODS

Compute β vector such that the **RSS function is minimized**

$$\operatorname{argmin}_{\beta} \text{RSS} = \|Y - \beta X\|^2$$

$$\operatorname{argmin}_{\beta} \text{RSS} = (Y - \beta X)^T(Y - \beta X)$$

In the Equation above, β is a quadratic function with $p+1$ parameters



IoT Sensor Data Forecasting

□ Linear Regression

$$\hat{y} = \hat{f}(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p$$

LINEAR LEAST SQUARE METHODS

$$\operatorname{argmin}_{\beta} RSS = (Y - \beta X)^T (Y - \beta X)$$

Differentiating with respect to β we get:

$$\frac{\partial RSS}{\partial \beta} = -2XT(Y - X\beta) = 0$$

↓

$$\beta = (X^T X)^{-1} X^T Y$$



IoT Sensor Data Forecasting

Generalizing the linear regression model ...

□ Polynomial Regression

$$\hat{y} = \hat{f}(\mathbf{x}) = \beta_0 + \beta_1 x_1 + \beta_2 x_2^2 + \cdots + \beta_p x_p^{d_p}$$

□ Basis Function

$$\hat{y} = \hat{f}(\mathbf{x}) = \beta_0 + \beta_1 b_1(x_1) + \beta_2 b_2(x_2) + \cdots + \beta_p b_p(x_p)$$

Functions $b_n(x_n)$ are **fixed** and **known** in advance



IoT Sensor Data Forecasting

□ How to assess the accuracy of the model?

Total **variance** of data:

$$\sum_{1}^n (y_i - \bar{y})^2 = \underbrace{\sum_{1}^n (\hat{y}_i - \bar{y})^2}_{\text{ESS}} + \underbrace{\sum_{1}^n (y_i - \hat{y}_i)^2}_{\text{RSS}}$$

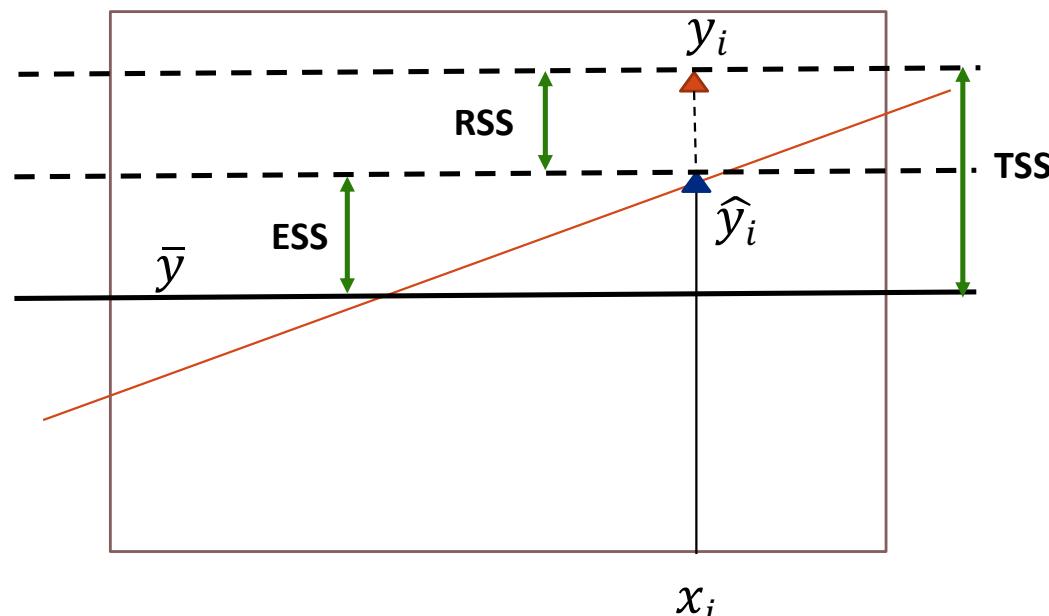


- **TSS** (Total Sum of Squares) → total variance in the response data
- **ESS** (Explained Sum of Squares) → variance explained by the model
- **RSS** (Residual Sum of Squares) → variance still present in the model



IoT Sensor Data Forecasting

- How to assess the accuracy of the model?



$$\begin{aligned} \text{TSS} &= \sum_{1}^n (y_i - \bar{y})^2 = \\ \text{RSS} &= \sum_{1}^n (y_i - \hat{y}_i)^2 + \\ \text{ESS} &= \sum_{1}^n (\hat{y}_i - \bar{y})^2 \end{aligned}$$



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

IoT Sensor Data Forecasting

□ How to assess the accuracy of the model?

$$R^2 = \frac{ESS}{TSS} = 1 - \frac{RSS}{TSS}$$

R^2 statistic:

- Lies between 0 and 1 → *Relative lack of fit*
- Measures the **proportion of variability in Y that can be explained by the regression**
- **Close to 1** → The model fits well the data
- **Close to 0** → Linear model is wrong or inherent error ε is high
- Acceptable value depends on the application scenario (e.g. marketing, biology, etc)
- It is a measure of the linear relationship between X and Y
(indeed, $R^2 = \text{Corr}(X, Y)^2$ for the simple case)



IoT Sensor Data Forecasting

STEP 1: LOAD THE DATASET

```
df = pd.read_csv("../tetuan.csv", parse_dates = ["DateTime"], index_col = "DateTime")
df["Time"] = np.arange(len(df.index))

X = pd.DataFrame(df["Time"])
y = pd.DataFrame(df["Temperature"])
```

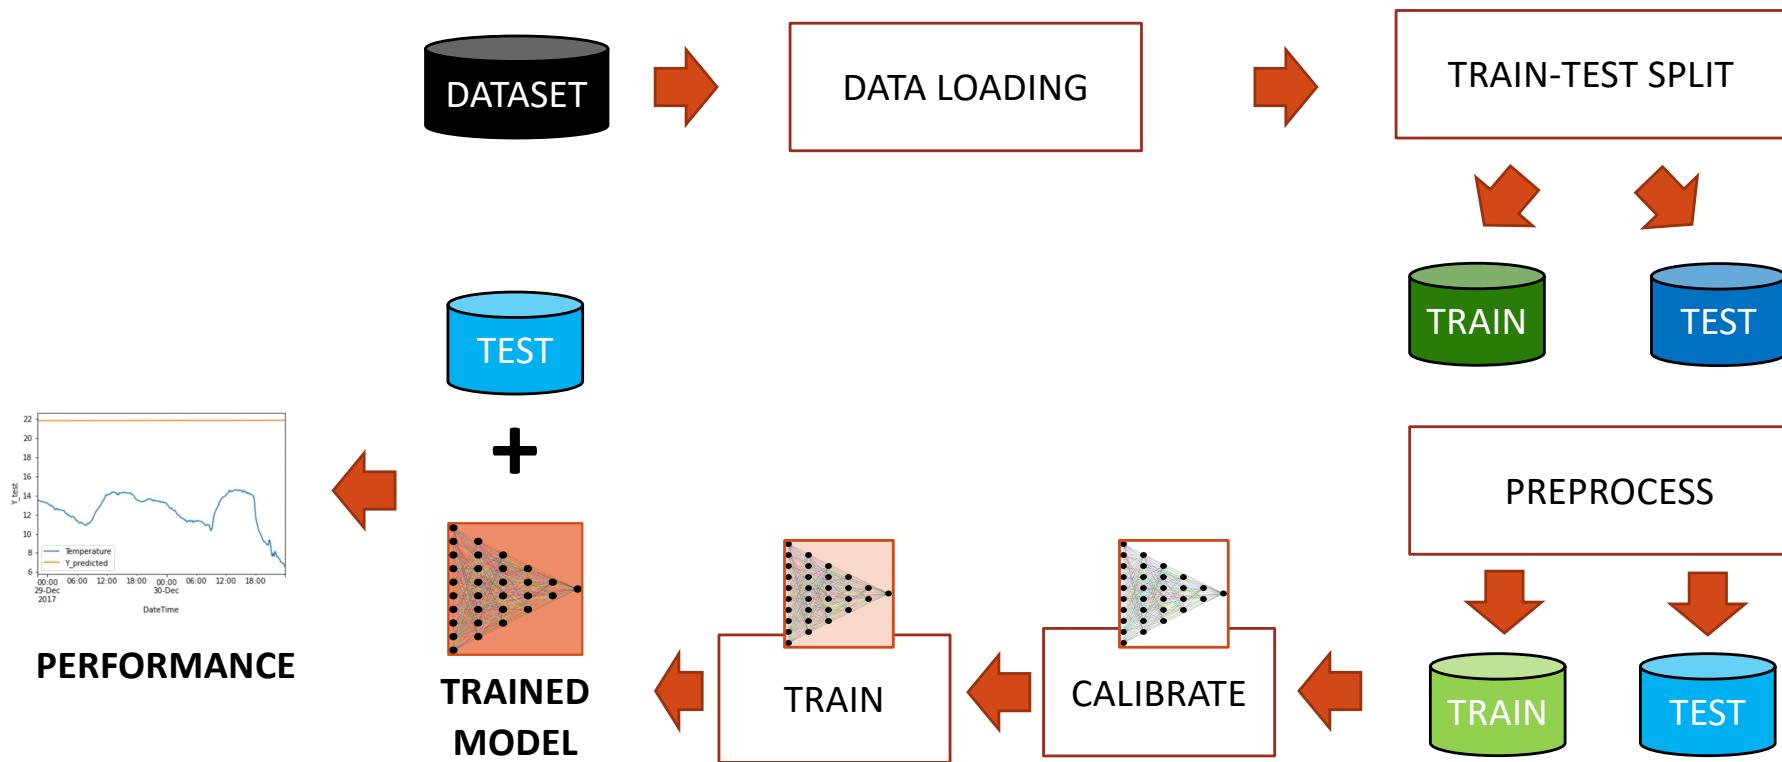
STEP 2: SPLIT THE DATASET

```
TOTAL_SAMPLES = len(df["Time"])
TEST_SIZE = 300

X_train = X.iloc[:TOTAL_SAMPLES - TEST_SIZE]
X_test = X.iloc[TOTAL_SAMPLES - TEST_SIZE:]
y_train = y.iloc[:TOTAL_SAMPLES - TEST_SIZE]
y_test = y.iloc[TOTAL_SAMPLES - TEST_SIZE:]
```



IoT Sensor Data Forecasting





IoT Sensor Data Forecasting

STEP 3: TRAIN THE MODEL

```
lreg = LinearRegression()  
lreg.fit(X_train, y_train)
```

STEP 4: PREDICT THE VALUES OF THE TEST

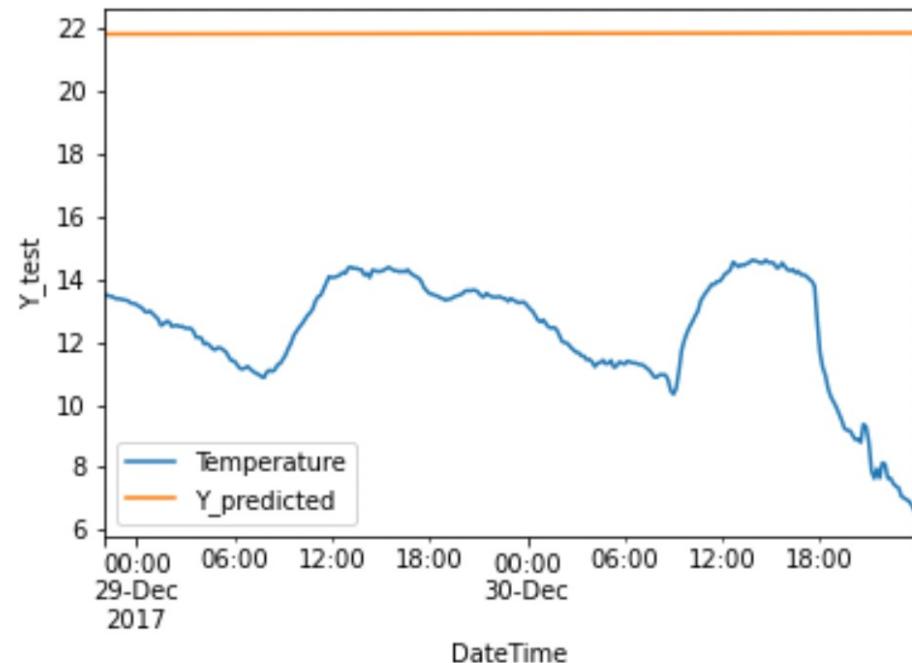
```
y_pred = lreg.predict(X_test)
```

STEP 5: COMPUTE THE STATS

```
mse=mean_squared_error(y_test, y_pred)  
mae=mean_absolute_error(y_test, y_pred)  
rsquare = r2_score(y_test, y_pred)
```



IoT Sensor Data Forecasting



```
Beta_0:  
[[0.00011495]]  
Beta_1:  
[15.8517469]  
Mean Square Error: 93.22180989659512  
Mean Absolute Error: 9.478330181662253
```



IoT Sensor Data Forecasting

AUTO-REGRESSION METHODS

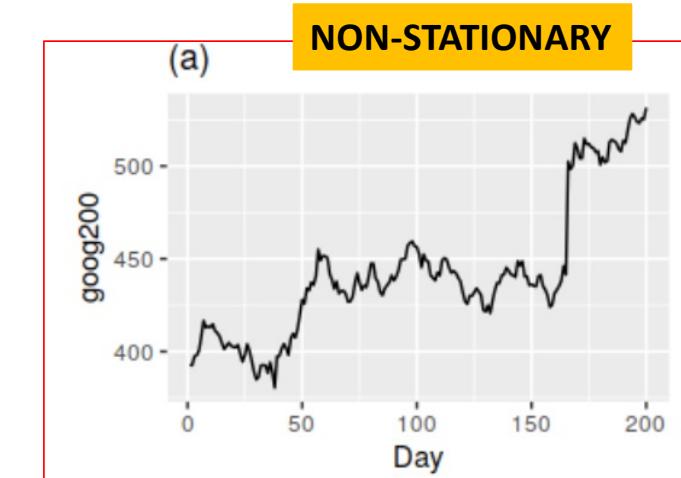
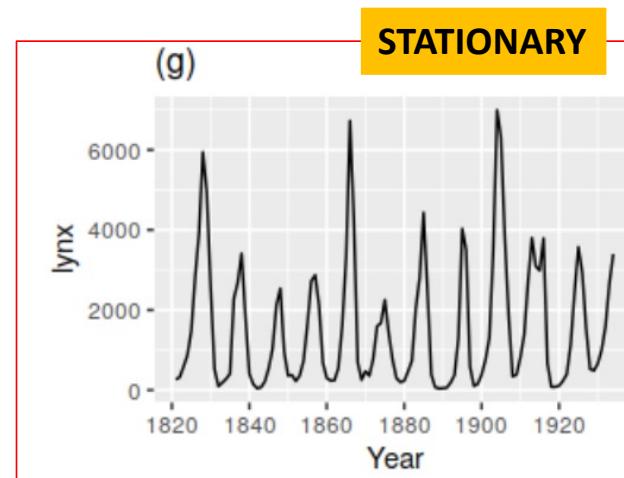
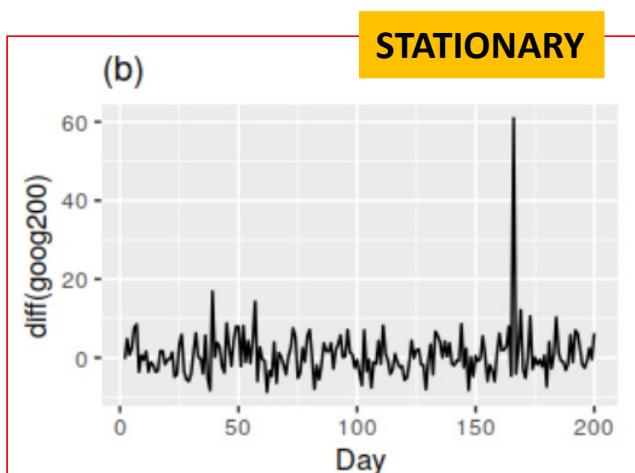


IoT Sensor Data Forecasting

✧ Time-series (**Weak**) Stationary

Source: <https://otexts.com/fpp2/stationarity.html>

- ✓ **Informal definition** → mean and variance of the time-series do not change over time.



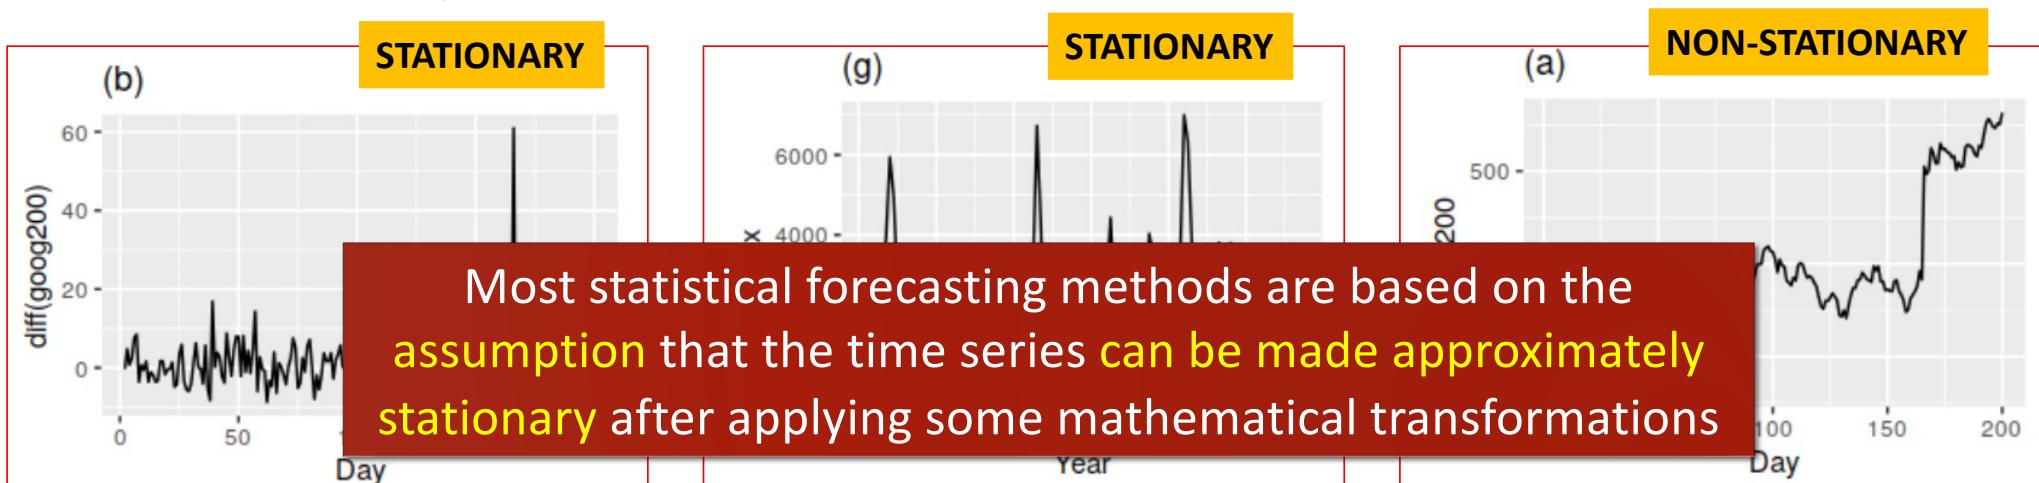


IoT Sensor Data Forecasting

✧ Time-series (**Weak**) Stationary

Source: <https://otexts.com/fpp2/stationarity.html>

- ✓ **Informal definition** → mean and variance of the time-series do not change over time.





IoT Sensor Data Forecasting

✧ **Lag** → delay between two points/timeseries

- X_t and X_{t+k} samples have lag k
- $\{X_0, X_1, \dots, X_t\}$ and $\{X_k, X_{k+1}, \dots, X_{k+t}\}$ are k -lagged time-series

STATIONARY TIME-SERIES

✧ **Autocorrelation Function (ACF)**

- ✧ Generally speaking, autocorrelation is the correlation of a signal with a delayed copy of itself (Wikipedia)
- ✧ It measures the degree of similarity between a given time series, and a lagged version of itself.

$$\rho_{X,Y} = \text{corr}(X, Y) = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y} = \frac{\mathbb{E}[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y}$$

Correlation function

$$r_k = \frac{\sum_{t=k+1}^n (y_t - \bar{y})(y_{t-k} - \bar{y})}{\sum_{t=1}^n (y_t - \bar{y})^2}$$

Autocorrelation



IoT Sensor Data Forecasting

✧ Moving Average (MA) Model

STATIONARY TIME-SERIES

❑ Q-th order moving average **model**:

$$Y_t = \mu + \varepsilon_t + \alpha_1 \cdot \varepsilon_{t-1} + \alpha_2 \cdot \varepsilon_{t-2} + \cdots + \alpha_q \cdot \varepsilon_{t-q}$$

- Weighted moving average of the past few forecast errors
- μ is the mean of the series
- ε_i are error terms (white noise terms, zero mean, iid)
- q is the order of MA model. How to choose it? ACF!

RULE. The autocorrelation function (ACF) of an MA(q) process is zero at lag $q + 1$ and greater.



IoT Sensor Data Forecasting

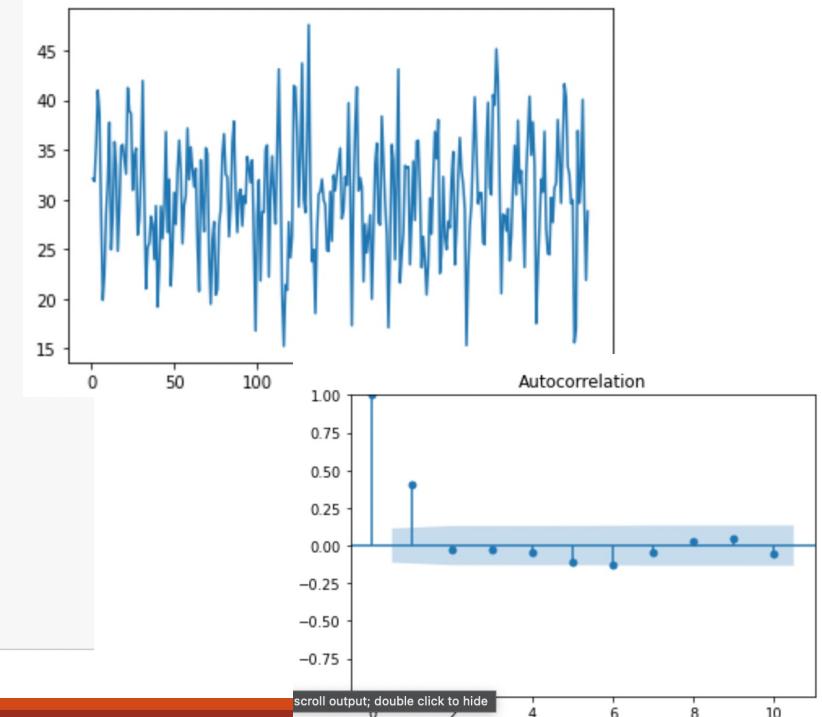
```
import numpy as np
from statsmodels.graphics import tsaplots

NUM_VALUES = 300
VARIANCE = 5
THETA = 0.5
MEAN = 30

e = np.random.normal(0, VARIANCE, NUM_VALUES)
y = []
x = range(1, NUM_VALUES)
for i in range(1, NUM_VALUES):
    val = MEAN + THETA * e[i-1] + e[i]
    y.append(val)
plt.plot(x,y)
fig = tsaplots.plot_acf(y, lags = 10)
```

MA(1) process

STATIONARY TIME-SERIES





IoT Sensor Data Forecasting

✧ Auto Regression (AR) Model

STATIONARY TIME-SERIES

- P-th order auto-regression model:

$$Y_t = c + \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \cdots + \phi_p Y_{t-p} + \varepsilon_t$$

- Value of the series at time t depends from past p values
- ε_t are white noise terms (zero mean, iid)
- p is the order of MA model. How to choose it?
 - ACF is not informative (too many parameters)
 - Use Partial Autocorrelation Function (**PACF**)



IoT Sensor Data Forecasting

✧ Partial Autocorrelation Function (PACF)

STATIONARY TIME-SERIES

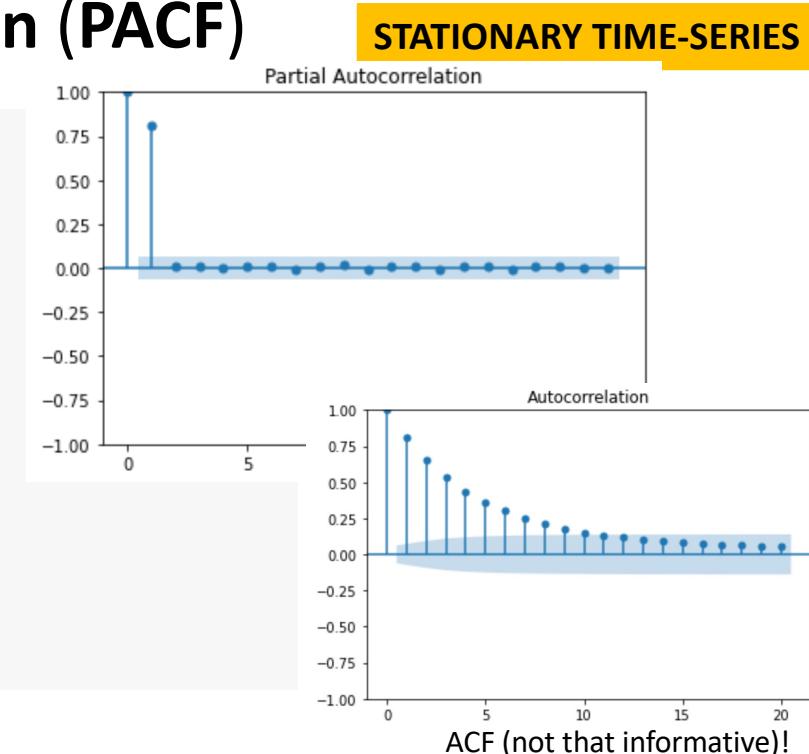
- Partial correlation of the time-series with itself at that lag given all the information between the points in time.
- IDEA: PACF between X_t and $X_{t+k} \rightarrow$ correlation between the two points removing the amount of correlation explained by $X_{t+1}, X_{t+2}, \dots, X_{t+k-1}$
- HOW to choose the p order for AR process?
 - The partial autocorrelation of an AR(p) process is zero at lag $p+1$ and greater.



IoT Sensor Data Forecasting

✧ Partial Autocorrelation Function (PACF)

```
1 from statsmodels.graphics import tsaplots
2 import matplotlib.pyplot as plt
3 import numpy as np
4 import math
5 import random
6
7
8 pval = 100
9
10 for i in range(0,100):
11     cval = 0.8 * pval + random.random()
12     x.append(cval)
13     pval = cval
14
15 fig = tsaplots.plot_pacf(x, lags=20)
16 plt.show()
```





IoT Sensor Data Forecasting

✧ ARMA(p,q) Model

STATIONARY TIME-SERIES

- Combination of p order Auto Regressive model and q order Moving Average model

$$Y_t = b + \varepsilon_t + \alpha_1 \cdot \varepsilon_{t-1} + \alpha_2 \cdot \varepsilon_{t-2} + \dots + \alpha_q \cdot \varepsilon_{t-q} + \\ \phi_1 \cdot Y_{t-1} + \phi_2 \cdot Y_{t-2} + \dots + \phi_{t-p} \cdot Y_{t-p}$$

- Estimate $p \rightarrow$ Use ACF Plot
- Estimate $q \rightarrow$ Use PACF Plot



IoT Sensor Data Forecasting

✧ How to decide whether the time-series is stationary?

□ Dickey-Fuller (ADF) test

- The test posits a null hypothesis that a unit root is present in a time series.
- Depending on the result of the test, the null hypothesis can be rejected for a specified significance level (series is stationary) or accepted (series is not stationary)

INTUITION: AR(1) Model

$$Y_t = c + \phi \cdot Y_{t-1} + \varepsilon_t$$

Null Hypothesis is true →
unit root is present →
 $\phi=1$ → mean not constant →
Series not stationary!

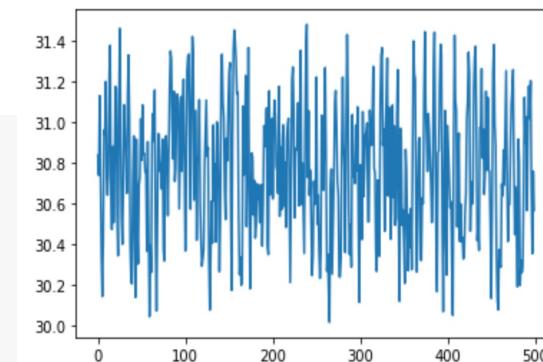


IoT Sensor Data Forecasting

✧ How to decide whether the time-series is stationary?

□ Dickey-Fuller (ADF) test

```
1 from statsmodels.tsa.stattools import adfuller  
2 import math  
3  
4 ALPHA = 0.5  
5 MEAN = 30  
6 NUM_SAMPLES = 500  
7 perror = 0  
8 y = []  
9 x = [ i for i in range(0,NUM_SAMPLES) ]  
10 for i in x:  
11     cerror = random.random()  
12     cval = cerror + MEAN + perror * ALPHA  
13     y.append(cval)  
14     perror = cerror  
15  
16 result = adfuller(y)  
17 print('ADF Statistic: %f' % result[0])  
18 print('p-value: %f' % result[1])
```



p – value is lower than confidence threshold (0.05) → Hypothesis rejected → **Stationary Series !** (it is AR(1) ...)

ADF Statistic: -9.016135
p-value: 0.000000

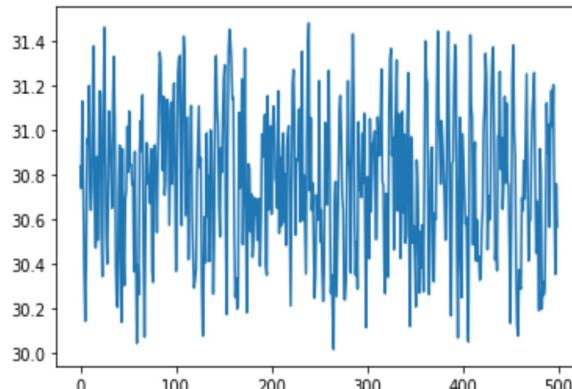


IoT Sensor Data Forecasting

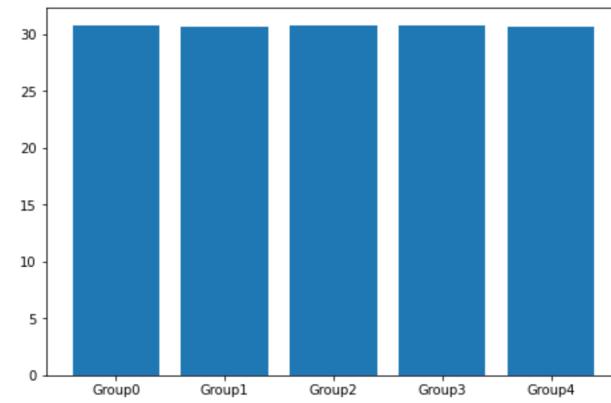
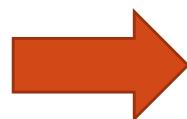
✧ How to decide whether the time-series is stationary?

□ Visual check

- Split the time series into two (or more) partitions and compare the mean variance of each group.

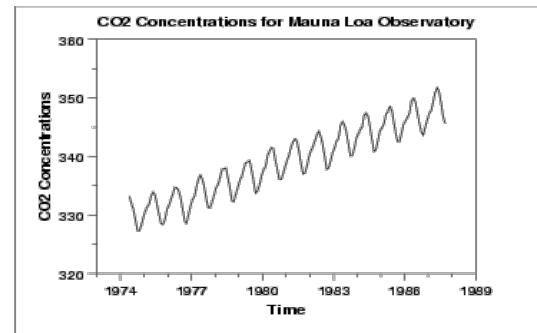
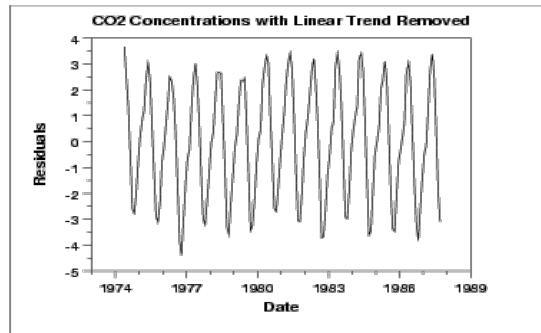


Mean per group
(group size:
100 samples)





IoT Sensor Data Forecasting



NON-STATIONARY TIME-SERIES

Source: <https://www.itl.nist.gov/div898/handbook/pmc/section4/pmc442.htm>

- Transforming a *non-stationary* time-series into a *stationary* time-series

Differencing technique → Stabilise the **mean** of a time series by removing changes over time

$$y'_t = y_t - y_{t-1}.$$

FIRST ORDER DIFFERENCING

$$\begin{aligned} y''_t &= y'_t - y'_{t-1} \\ &= (y_t - y_{t-1}) - (y_{t-1} - y_{t-2}) \\ &= y_t - 2y_{t-1} + y_{t-2}. \end{aligned}$$

SECOND ORDER DIFFERENCING



IoT Sensor Data Forecasting

NON-STATIONARY TIME-SERIES

□ ARIMA(p,q,d) Autoregressive Integrated Moving Average Model

- **p** → number of **autoregressive** terms (AR order)
- **d** → number of **nonseasonal differences** (differencing order)
- **q** → number of **moving-average** terms (MA order)

$$y'_t = c + \phi_1 y'_{t-1} + \cdots + \phi_p y'_{t-p} + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t;$$

y'_t is the **d-differenced** value at time t



IoT Sensor Data Forecasting

```
1 #Step 1: Load the dataset
2
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6
7 df = pd.read_csv('airlines_data.csv')
```

PUTTING ALL TOGETHER ...

```
1 #Step 2: Test-train split
2 from datetime import datetime
3
4 df['Month']=pd.to_datetime(df['Month'],infer_datetime_format=True)
5 df = df.set_index('Month',inplace=False)
6
7 nrows = (len(df.values))
8 splitPoint = int(nrows * 0.80)
9
10 train = df['Passengers'][:splitPoint]
11 test = df['Passengers'][splitPoint:]
```

LOAD THE DATASET
AND TRAIN-TEST SPLIT ...



IoT Sensor Data Forecasting

```
1 #Step 3: Check whether the time-series is stationary through the Dickey-Fuller test
2
3 from statsmodels.tsa.stattools import adfuller
4
5 result = adfuller(train)
6 print('ADF Statistic: %f' % result[0])
7 print('p-value: %f' % result[1])
```

ADF Statistic: -0.356889
p-value: 0.917052

IS IT STATIONARY? ... NO!

```
1 #Step 4: Find the differencing parameter (d)
2 from statsmodels.graphics import tsaplots
3
4 trainNew = train.diff().dropna()
5
6 result = adfuller(trainNew)
7 print('ADF Statistic: %f' % result[0])
8 print('p-value: %f' % result[1])
```

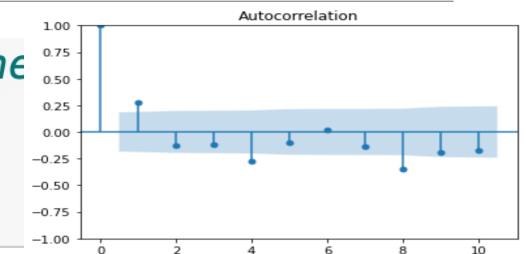
ADF Statistic: -2.539635
p-value: 0.106126

REPEAT THE TEST AFTER 1ST-ORDER DIFFERENCING
IS IT STATIONARY? YES (threshold=0.1)



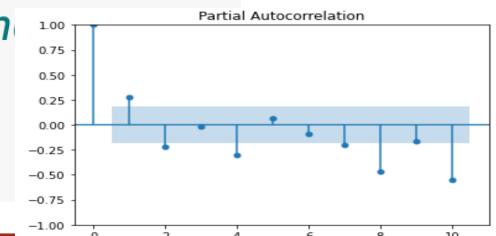
IoT Sensor Data Forecasting

```
1 #Step 4: Find the order of the MA model through the  
2  
3 fig = tsaplots.plot_acf(trainNew, lags=10)  
4 plt.show()
```



FIND p ORDER FOR THE MA MODEL
COMPUTE ACF → p = 1

```
1 #Step 5: Find the order of the AR model through the  
2  
3 fig = tsaplots.plot_pacf(trainNew, lags=10)  
4 plt.show()
```



FIND q ORDER FOR THE AR MODEL
COMPUTE PACF → q = 1



IoT Sensor Data Forecasting

```
1 #Step 6: Apply ARIMA model
2
3 from statsmodels.tsa.arima.model import ARIMA
4
5 history = [x for x in train]
6 predictions = list()
7
8 for t in range(len(test)):
9     model = ARIMA(history, order=(1,1,1))
10    model_fit = model.fit()
11    output = model_fit.forecast()
12    yest = output[0]
13    predictions.append(yest)
14    obs = test[t]
15    history.append(obs)
16    print('predicted=%f, expected=%f' % (yest, obs))
```

APPLY ARIMA(p,q,d)

FORECAST ALL SAMPLES
IN THE TEST DATASET



IoT Sensor Data Forecasting

✧ Mean Absolute Error (MAE)

- Measures the **total error in forecast** (without sign)

$$\text{MAD} = \frac{\sum |\text{actual} - \text{forecast}|}{n}$$

✧ Mean Square Error (MSE)

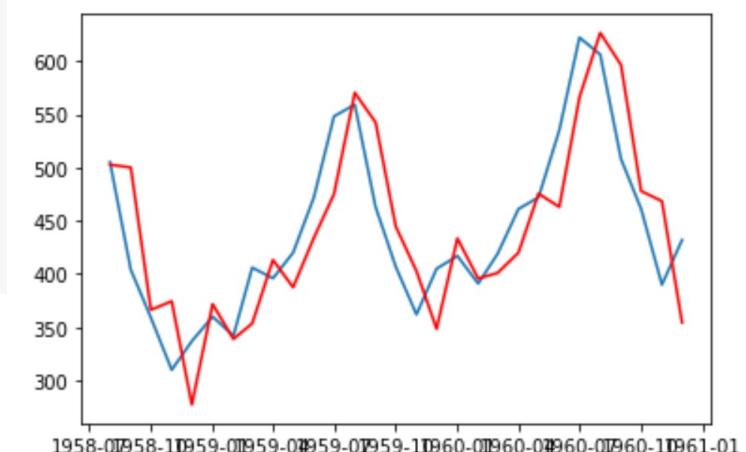
- Similar to MAD, but gives **higher penalty** to larger errors

$$\text{MSE} = \frac{\sum (\text{actual} - \text{forecast})^2}{n}$$



IoT Sensor Data Forecasting

```
1 import math
2 from sklearn.metrics import mean_squared_error
3
4 rmse = math.sqrt(mean_squared_error(test, predictions))
5 print('Test RMSE: %.3f' % rmse)
6
7 df2 = pd.DataFrame(predictions)
8 df2.set_index(test.index, inplace=True)
9
10 plt.plot(test)
11 plt.plot(df2, color='red')
12 plt.show()
```





IoT Sensor Data Forecasting

BEYOND ARIMA 1/2

- ❑ Python module `pmdarima` allows to compute the optimal ARIMA parameters (p, d, and q values). HOW?
- ❑ Grid search technique (testing parameters in a range)

```
from pmdarima.arima import auto_arima

model = auto_arima(trainset, start_p=1, start_q=1, test='adf', max_p=5,
max_q=5, m=1, d=1, seasonal=False, start_P=0, D=None, trace=True,
error_action='ignore', suppress_warnings=True, stepwise=True)

prediction, confint = model.predict(n_periods=TEST_SIZE,
return_conf_int=True)
```



IoT Sensor Data Forecasting

BEYOND ARIMA 2/2

- ❑ HOW to handle with the seasonality?
- ❑ **SARIMA (Seasonal ARIMA)**

$\text{SARIMA}(p, d, q)(P, D, Q)_m$

Four additional components:

- P: **Seasonal** autoregressive order.
- D: **Seasonal** difference order.
- Q: **Seasonal** moving average order.
- m: The number of time steps for a single seasonal period.



IoT Sensor Data Forecasting

BEYOND ARIMA 2/2

- ❑ HOW to handle with the seasonality?
- ❑ **SARIMA (Seasonal ARIMA)**

$\text{SARIMA}(p, d, q)(P, D, Q)_m$

```
from pmdarima.arima import auto_arima  
  
arima_model_auto = auto_arima(y_train, seasonal=True, m=12, trace=True)
```



IoT Sensor Data Forecasting

PROPHET LIBRARY FOR TIME-SERIES FORECASTING

<https://facebook.github.io/prophet/>

CONNECTING PROPHET TO INFLUXDB

<https://www.influxdata.com/blog/forecasting-with-fb-prophet-and-influxdb/>

- Additive Regression Model
- It works best with time-series with strong **seasonal** effects and several seasons of data
- It estimates yearly, weakly and daily seasonality plus the holiday effect
- **Automatic** method (no pre-preprocessing, no parameter tuning)



IoT Sensor Data Forecasting

```
from datetime import datetime

df['Month'] = pd.DatetimeIndex(df['Month'])
df = df.rename(columns={'Month': 'ds', 'Passengers': 'y'})
df.set_index('ds')
```

```
1 from fbprophet import Prophet
2
3 model = Prophet(interval_width=0.95)
4 model.fit(df)
5 future_dates = model.make_future_dataframe(periods=36, freq='MS')
6 forecast = model.predict(future_dates)
7 model.plot(forecast, uncertainty=True)
8 model.plot_components(forecast)
```



IoT Sensor Data Forecasting

